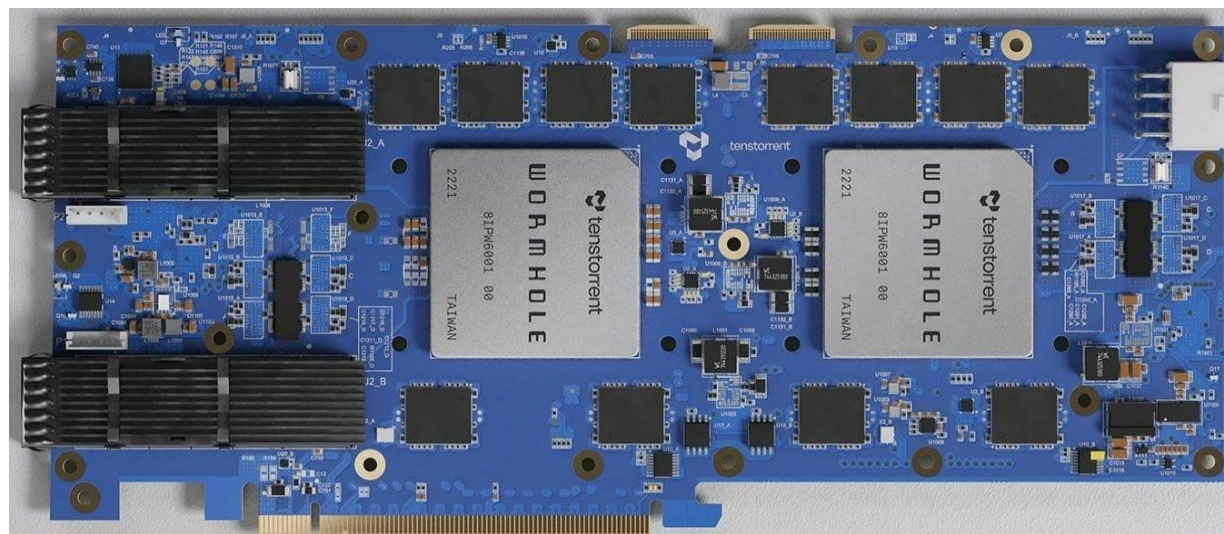
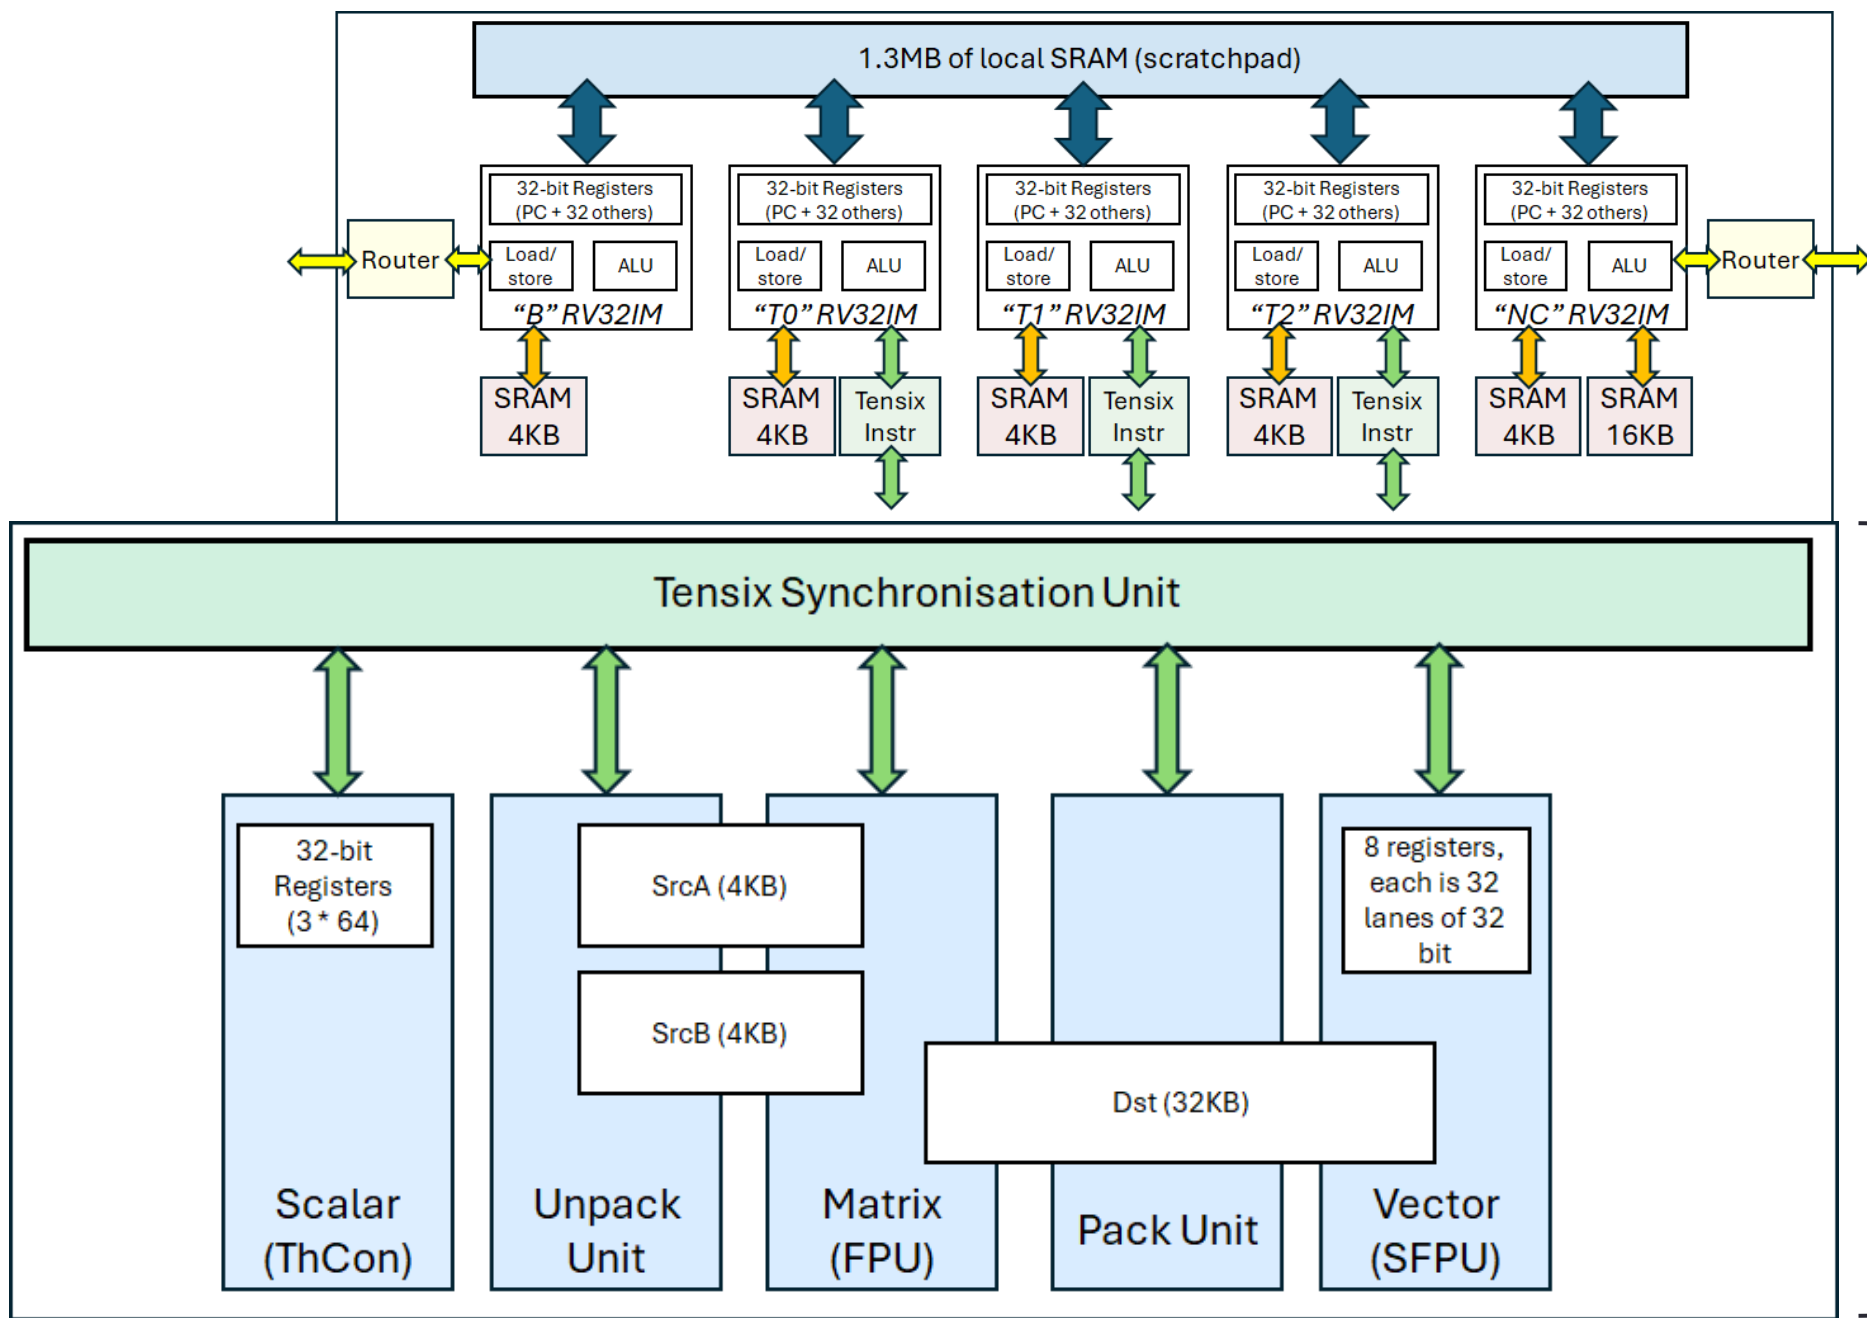


Overview of TT-Metalium SDK: Compute

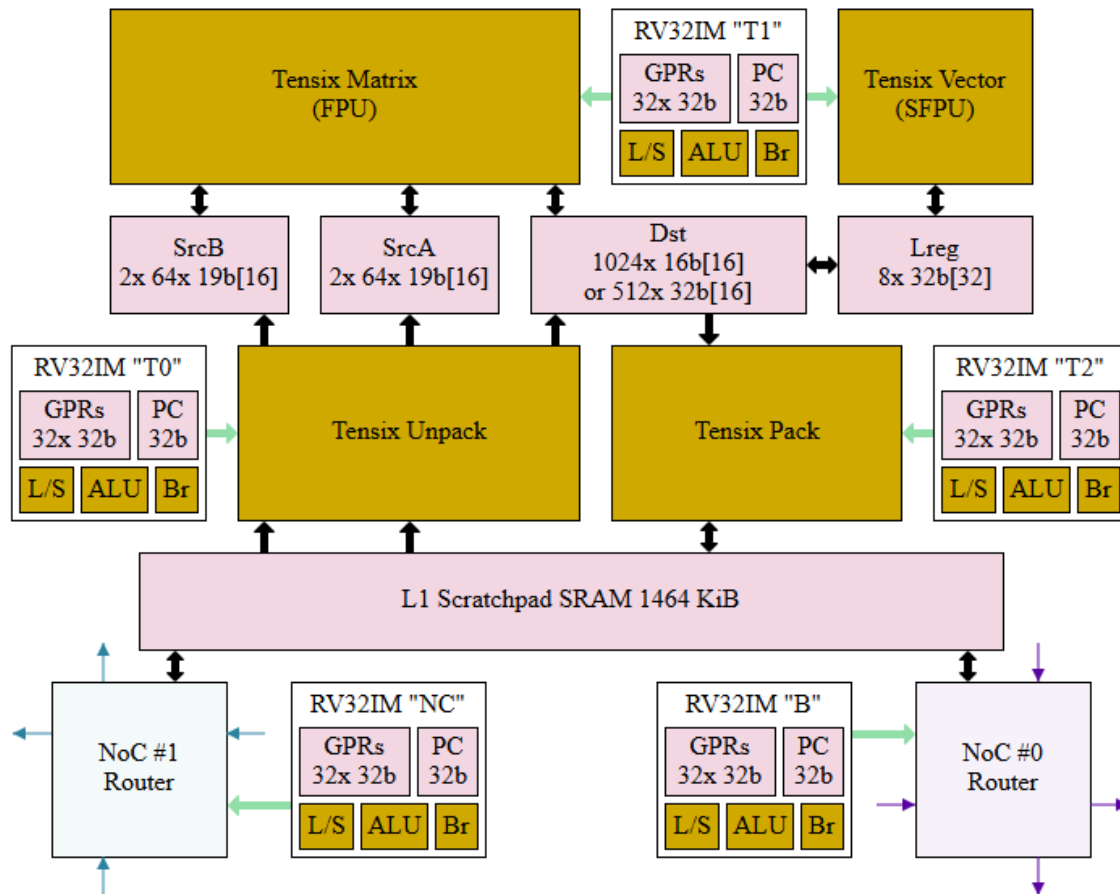




The compute engine has a scalar, matrix and vector unit

- SrcA and SrcB are input registers
- Dst is an output register (and input register for Vector unit too)

A more accurate diagram....




- The matrix unit can perform up to 4.096 TFLOP/s
- The vector unit can perform 32 FP32 maths operations per cycle
- This is FP32(ish)
 - The matrix unit makes several sacrifices to IEEE compliance to achieve this performance
 - Maximum of 19 bits for each element, so supports a maximum of TF32
 - The vector unit provides full FP32 support, so for HPC workloads this is probably the one we would use
 - However, performance is slower

From Pete Cawley's blog at <https://www.corsix.org/content/tt-wh-part7>

Data supported by the matrix unit (FPU)

- Inputs to the matrix unit (srcA and srcB) are two banks of 64 rows by 16 columns of 19-bit data
 - 1024 elements computed with by the FPU, the majority of compute FPU operations have a latency of five cycles but are pipelined and can be issued each cycle
 - Matrix multiplication can deliver up to 4.096 TFLOP/s and element wide operations 0.256 TFLOP/s



<u>Dst</u> data type	+=	SrcB data type	@	SrcA data type
8x16 matrix of either FP32 or BF16	+=	8x16 matrix of either TF32 or BF16	@	16x16 matrix of either TF32 (†) or BF16
8x16 matrix of either FP32 or FP16	+=	8x16 matrix of FP16	@	16x16 matrix of FP16 (†)
8x16 matrix of integer "32"	+=	8x16 matrix of integer "8"	@	16x16 matrix of integer "8" (‡)

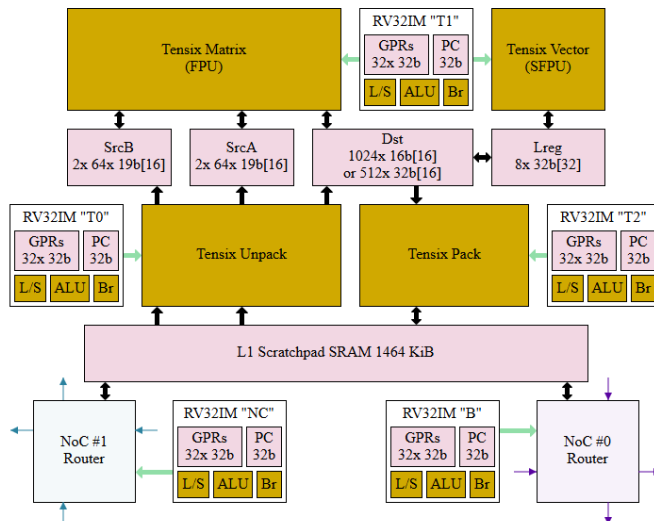
From <https://github.com/tenstorrent/tt-isa-documentation/blob/main/WormholeB0/TensixTile/TensixCoproprocessor/SrcASrcB.md>

- Whilst the FPU provides the potential for performance, the data types that are supported are rather limited

Issuing compute operations to matrix unit

To get input data in

- Wait for two CBs (LHS and RHS) to be available via *cb_wait_front* API call



To compute


- Acquire exclusive compute lock on DST (target) registers using *tile_regs_acquire* API call
- Issue corresponding matrix API call such as *add_tiles*, *sub_tiles*, *mul_tiles* with CB index as input
- Release exclusive compute lock on DST (target) registers using *tile_regs_commit* API call

To get results out

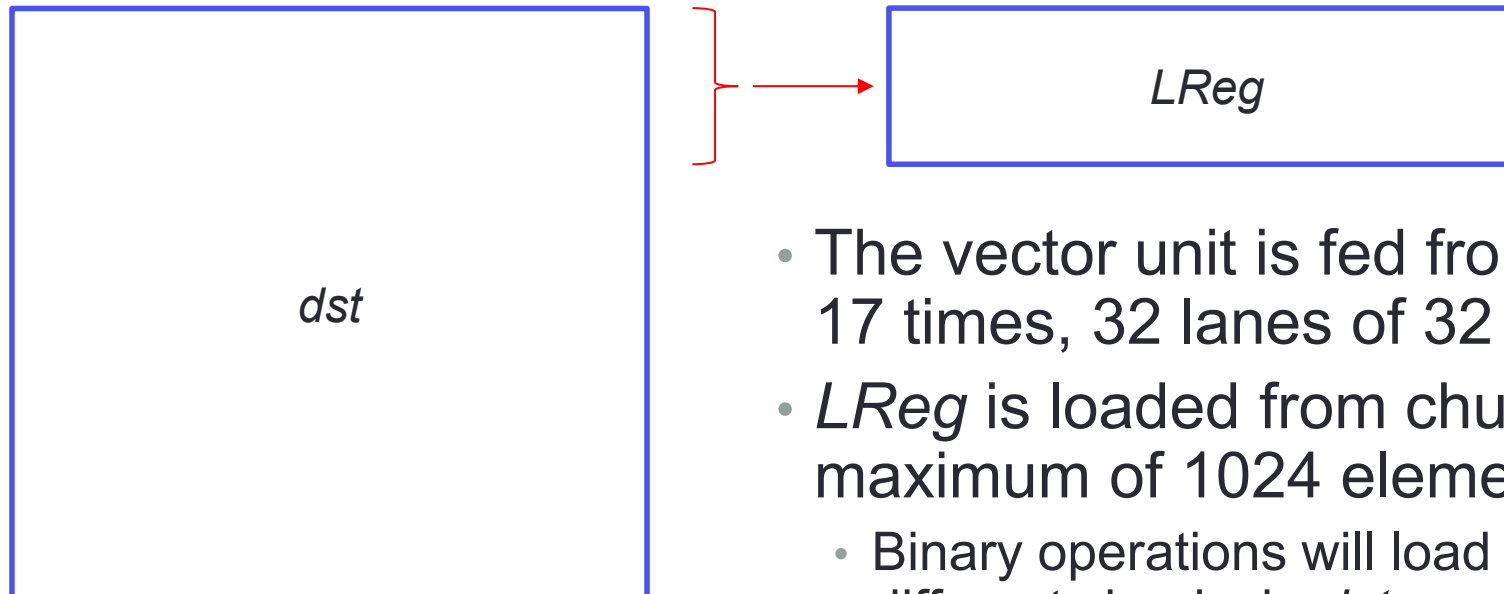
- Acquire exclusive pack lock on DST (target) registers using *tile_regs_wait* API call
- Copy results from dst register to target CB via *pack_tile* API call
- Release exclusive pack lock on DST (target) registers using *tile_regs_release* API call

Data supported by the vector unit (SFPU)

- Inputs to the vector unit are from the *dst* register
 - 1024 rows of 16 columns of 16-bit data, or 512 rows of 16 columns of 32-bit data
 - BF16, FP16, FP32, int8, int16 and int32 are supported

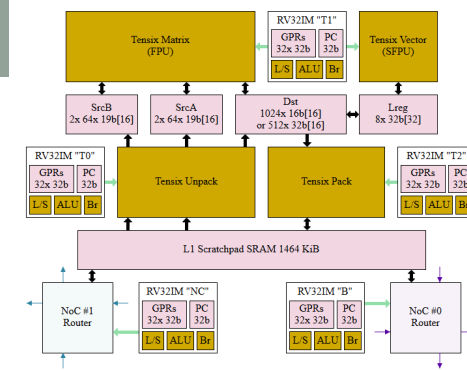


- The vector unit has 32 lanes, each of 32 bit



- The vector unit is fed from *LReg*, which contains 17 times, 32 lanes of 32 bit
- LReg* is loaded from chunks of *dst*, from a maximum of 1024 elements
 - Binary operations will load different lanes of *LReg* from different chunks in *dst*

Issuing compute operations to vector unit



To get input data in

- Wait for two CBs (LHS and RHS) to be available via *cb_wait_front* API call
- Acquire exclusive pack lock on DST (target) registers using *tile_regs_wait* API call
- Copy both input tiles into *dst* register using segment index
- Release exclusive pack lock on DST (target) registers using *tile_regs_release* API call

To compute

- Acquire exclusive compute lock on DST (target) registers using *tile_regs_acquire* API call
- Issue corresponding vector API call such as *add_binary_tile*, *sub_binary_tile* with segment index determining inputs (first input overwritten with results)
- Release exclusive compute lock on DST (target) registers using *tile_regs_commit* API call

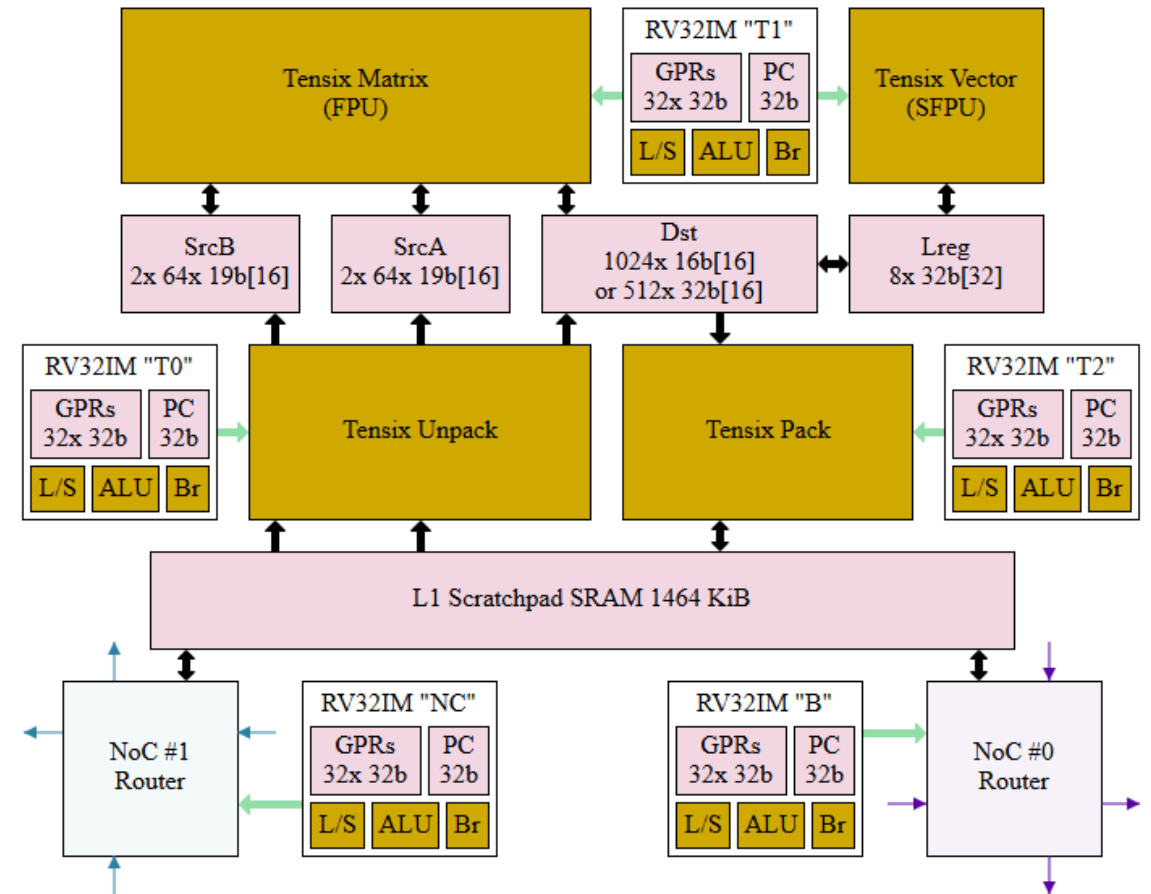
To get results out

- Acquire exclusive pack lock on DST (target) registers using *tile_regs_wait* API call
- Copy results from dst register to target CB via *pack_tile* API call
- Release exclusive pack lock on DST (target) registers using *tile_regs_release* API call

Vector unit uses the dst registers for both inputs and output

The key points

- Need to initialise with the data type
 - And reinitialise if change this
- Inputs are CBs and the output is a CB
- *dst* register is split into 16 segments
 - Matters more when using the vector unit
- Need to acquire locks on the *dst* register as this coordinates instructions from the pack, compute and unpack RISC-V cores



Initialisation

- The compute engine must be initialised, taking the input and output circular buffers as arguments
 - This configures the unpacker, packer, and FPU for the specific operation being performed.
 - Re-initialization is not required for repeated operations with the same source, destination, and data type parameters.
- In practical four (FPU) you will see we use *binary_op_init_common* (with input and output CBs) and *add_tiles_init* (with input CBs)
- In practical five (SFPU) we use *init_sfpu* (with an input and output CB) and *add_int_tile_init* (with no arguments)
 - *init_sfpu* sets up the packer and unpackers, but there are limits (assumes all CBs the same type and doesn't support FP32) and *copy_tile_init* can be used instead to get round these

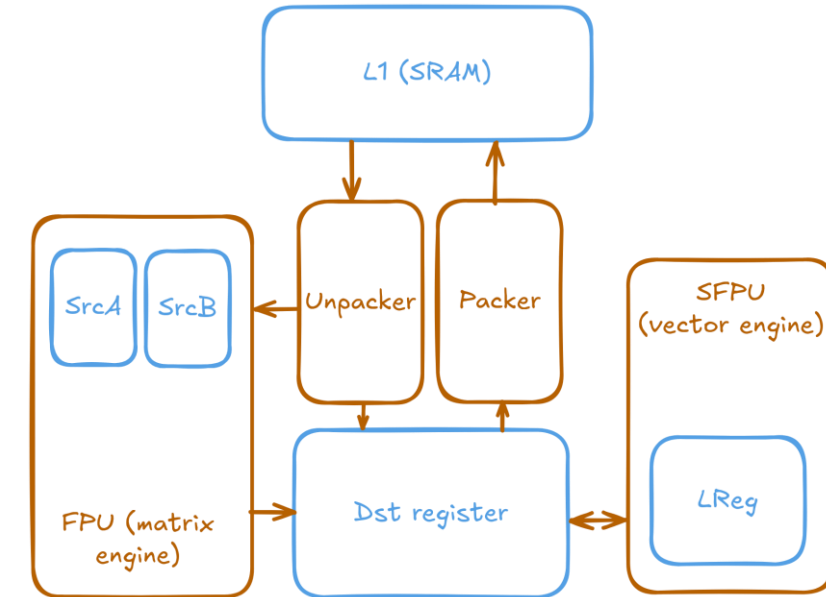


Image from https://docs.tenstorrent.com/tt-metal/latest/tt-metalium/tt_metal/advanced_topics

Most common maths calls

Matrix unit

- add_tiles
- sub_tiles
- mul_tiles
- matmul_tiles
- reduce_tile
- transpose_wh_tile

This also explains why the vector unit consumes from dst, as a common ML use-case is to execute with the matrix unit and then run another operation on results via the vector unit

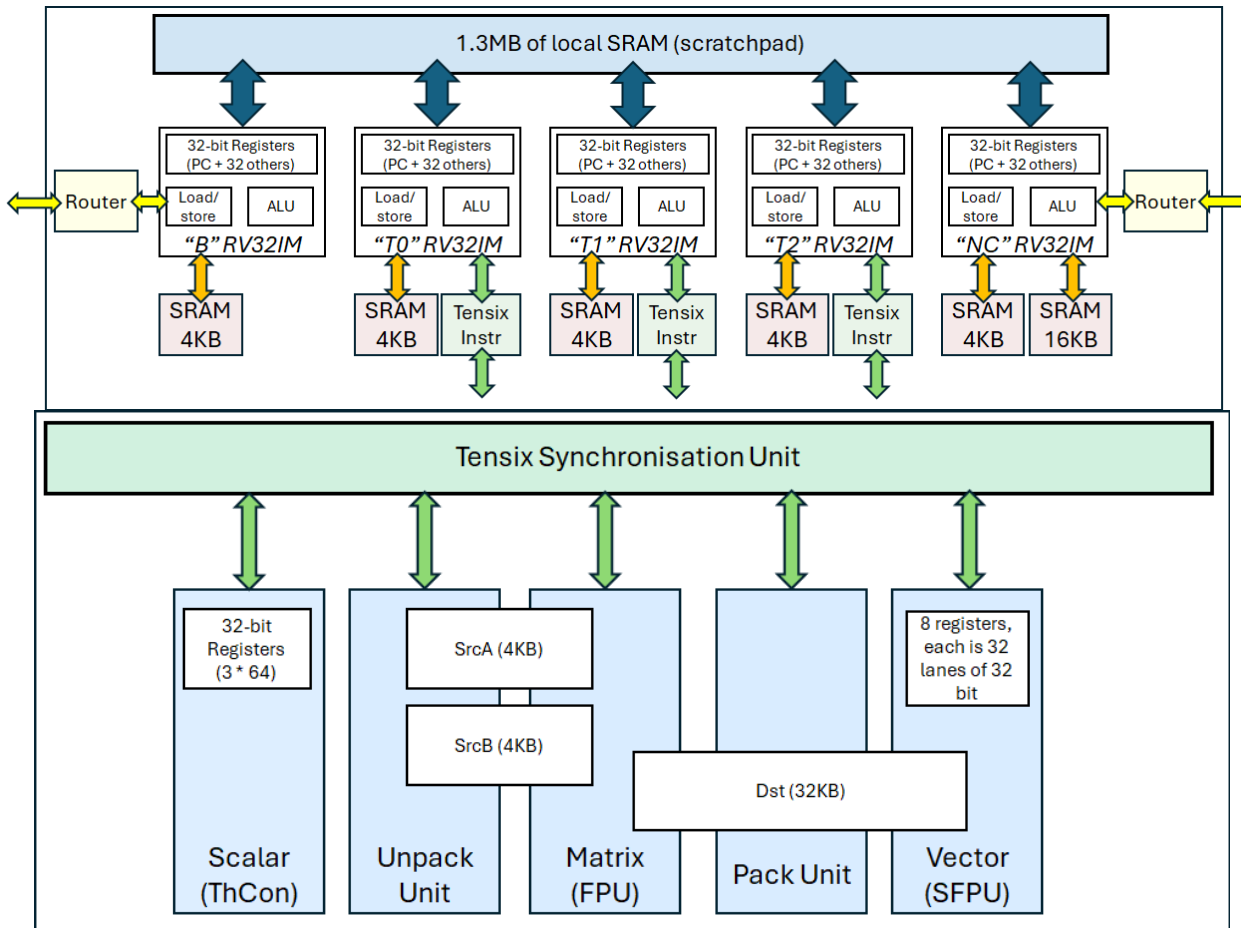
Vector unit

- add_binary_tile
 - sub_binary_tile
 - mul_binary_tile
 - abs_tile
 - exp_tile
 - lsinf_tile
 - lsfinitie_tile
 - lsnan_tile
 - sqrt_tile
 - square_time
 - tan_tile
 - sin_tile
 - cos_tile
 - atan_tile
 - acos_tile
 - asin_tile
- And integer variants*

- ltz_tile
- eqz_tile
- lez_tile
- gtz_tile
- neq_tile
- gez_tile
- unary_ne_tile
- unary_gt_tile
- unary_lt_tile
- unary_max_tile
- unary_min_tile

This column operations are for comparison

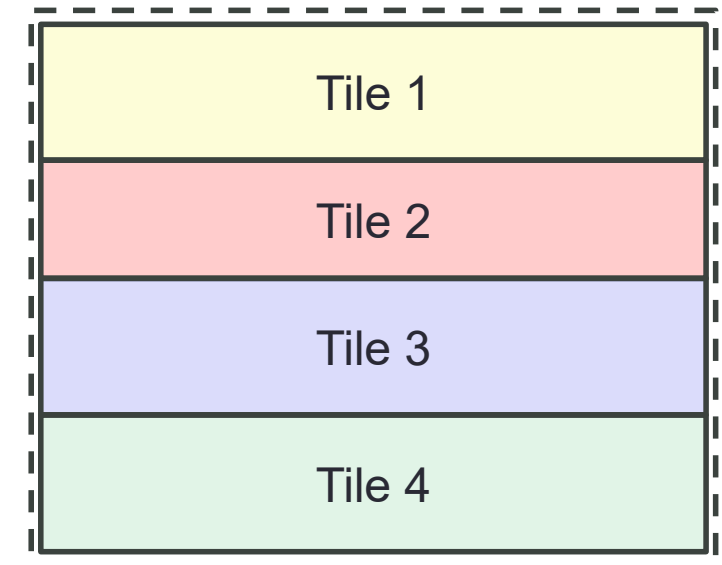
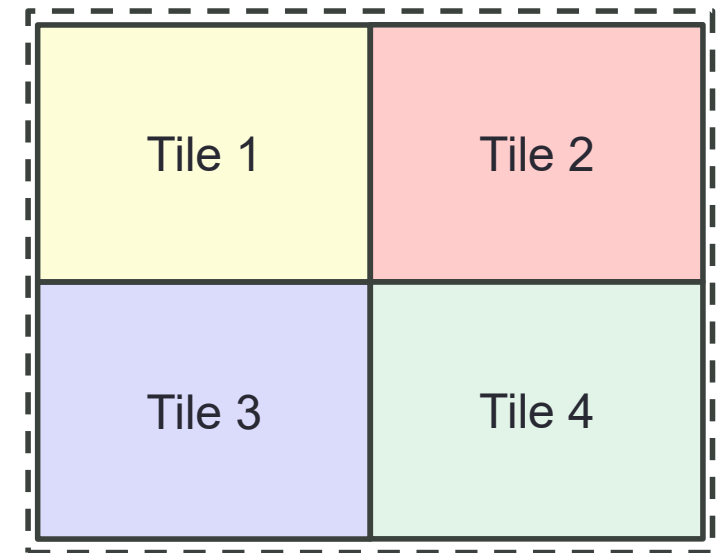
Three compute cores...



- Each Tensix unit has three RISC-V baby cores for compute
 - Unpacker drives the unpack unit
 - Maths drives FPU, SFPU, ThCon
 - Packer drives the pack unit
- Programmer's compute kernel is launched on all three cores which execute it concurrently
 - In the Metalium API there are explicit sections for different cores, where one (or more) cores will execute some code and additional synchronisation
 - But we don't really need to worry about this, however, it explains why there are locks on the dst register to avoid conflict between the math and pack cores

Tiling data to drive compute

- We have talked about bringing the FPU into play but the registers are only of a certain size
 - srcA and srcB contain a maximum of 1024 elements and similar if you use the SFPU
- Therefore need to tile data across chunks
 - Tenstorrent use the terminology *tile* due to the architecture being designed for matrix multiplications, and chunk would be better as a tile can be 1D
- Practical three will explore how to do this, before using the matrix multiplication engine in practical four and vector unit in practical five to perform the compute



What now: Practicals 3,4 & 5 and more information

- We are now going to move onto looking at practicals 3, 4, and 5
 - Practical three chunks up data (into tiles) and operates on each chunk in turn. This is in preparation to use the FPU/SFPU as they have a maximum of 1024 elements at a time
 - Practical four explores using the FPU (matrix unit) to undertake element wise addition. We sacrifice accuracy (can only use int8) but get performance.
 - Practical five explores using the SFPU (vector unit) to undertake element wise addition. We sacrifice performance but get accuracy (can use int32).
- There is API based documentation at https://docs.tenstorrent.com/tt-metal/latest/tt-metalium/tt_metal/apis/index.html
 - But this is somewhat incomplete, the API header files at https://github.com/tenstorrent/tt-metal/tree/main/tt_metal/include/compute_kernel_api tend to be more useful
- Detailed architecture documentation can be found at <https://github.com/tenstorrent/tt-isa-documentation/tree/main/WormholeB0>